

CS4487 - Machine Learning

Lecture 4b - Non-linear Classifiers

Dr. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

Outline

1. Nonlinear classifiers
2. Kernel trick and kernel SVM
3. **Ensemble Methods - Boosting, Random Forests**
4. Classification Summary

Ensemble Classifiers


- *Why trust only one expert?*
 - In real life, we may consult several experts, or go with the "wisdom of the crowd"
 - In machine learning, *why trust only one classifier?*
- Ensemble methods aim to combine multiple classifiers together to form a better classifier.
- Examples:
 - **boosting** - training multiple classifiers, each focusing on errors made by previous classifiers.
 - **bagging** - training multiple classifiers from random selection of training data

AdaBoost - Adaptive Boosting

- Base classifier is a "weak learner"
 - A simple classifier that can be slightly better than random chance (>50%)
 - Example: *decision stump classifier*
 - check if feature value is above (or below) a threshold.
 - $y = f(x) = \begin{cases} +1, & x_j \geq T \\ -1, & x_j < T \end{cases}$

```
In [4]: wlf1g
```

```
Out[4]:
```




- **Idea:** train weak classifiers sequentially
- In each iteration,
 - Pick a weak learner $f_i(\mathbf{x})$ that best carves out the input space.
 - The weak learner should focus on data that is misclassified.
 - Apply weights to each sample in the training data.
 - Higher weights give more priority to difficult samples.

Iteration 1

- Initially, weights for all training samples are equal: $w_i = 1/N$
 - Pick best weak learner.

```
In [7]: plt[1]
```

```
Out[7]:
```




Iteration 2

- points are re-weighted based on the current classification result:
 - increase weights of samples that are misclassified: $w_i = w_i e^{\alpha}$
 - decrease weights of correctly classified samples: $w_i = w_i e^{-\alpha}$
 - $\alpha = 0.5 \log \frac{1-err}{err}$ is based on the current classifier error.
 - (larger circles indicates higher weight)
- using the weighted data, train another weak learner $f_2(\mathbf{x})$.
- the classifier function is the weighted sum of weak learners
 - $f(\mathbf{x}) = \sum_{t=1}^D \alpha_t f_t(\mathbf{x})$

```
In [8]: plt[2]
```


Out[8]:



Keep iterating...


```
In [9]: plt[3]
```

Out[9]:




```
In [10]: plt[4]
```

Out[10]:



```
In [11]: plt[5]
```


Out[11]:



- After many iterations...

```
In [13]: adafig
```

```
Out[13]:
```




Adaboost loss function

- exponential loss
 - $L(z_i) = e^{-z_i}$
 - $z_i = y_i f(\mathbf{x}_i)$
 - very sensitive to misclassified outliers.

```
In [15]: lossfig
```

```
Out[15]:
```




Example on Iris data

- Too many weak-learners and AdaBoost curves out space for the outliers.

```
In [19]: irisfig
```

```
Out[19]:
```



- use cross-validation to select number of weak learners.

```
In [20]: # setup the list of parameters to try
paramgrid = {'n_estimators': array([1, 2, 3, 5, 10, 15, 20, 25, 50, 100, 200, 500, 1000]) }
print(paramgrid)


# setup the cross-validation object
# (NOTE: using parallelization in GridSearchCV, not in AdaBoost)
adacv = model_selection.GridSearchCV(ensemble.AdaBoostClassifier(random_state=4487),
                                     paramgrid, cv=5, n_jobs=-1)

# run cross-validation (train for each split)
adacv.fit(trainX, trainY);

print("best params:", adacv.best_params_)

{'n_estimators': array([ 1,  2,  3,  5, 10, 15, 20, 25, 50, 100, 200, 500, 1000])}
best params: {'n_estimators': 2}
```

```
In [22]: (avgscores, pnames, bestind) = extract_grid_scores(adacv, paramgrid)
paramfig = plt.figure()
plt.semilogx(paramgrid['n_estimators'], avgscores, 'o-')
plt.grid(True)
plt.ylabel('accuracy'); plt.xlabel('number of weak learners')
plt.title('accuracy for different number of weak learners')
plt.show()
```




```
In [23]: # predict from the model
predY = adacv.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)
```


test accuracy = 0.82


```
In [25]: ifig2
```

```
Out[25]:
```



- Boosting can do feature selection
 - each decision stump classifier looks at one feature
- One of the original face detection methods (Viola-Jones) used Boosting.
 - extract a lot of image features from the face
 - during training, Boosting learns which ones are the most useful.



AdaBoost Summary

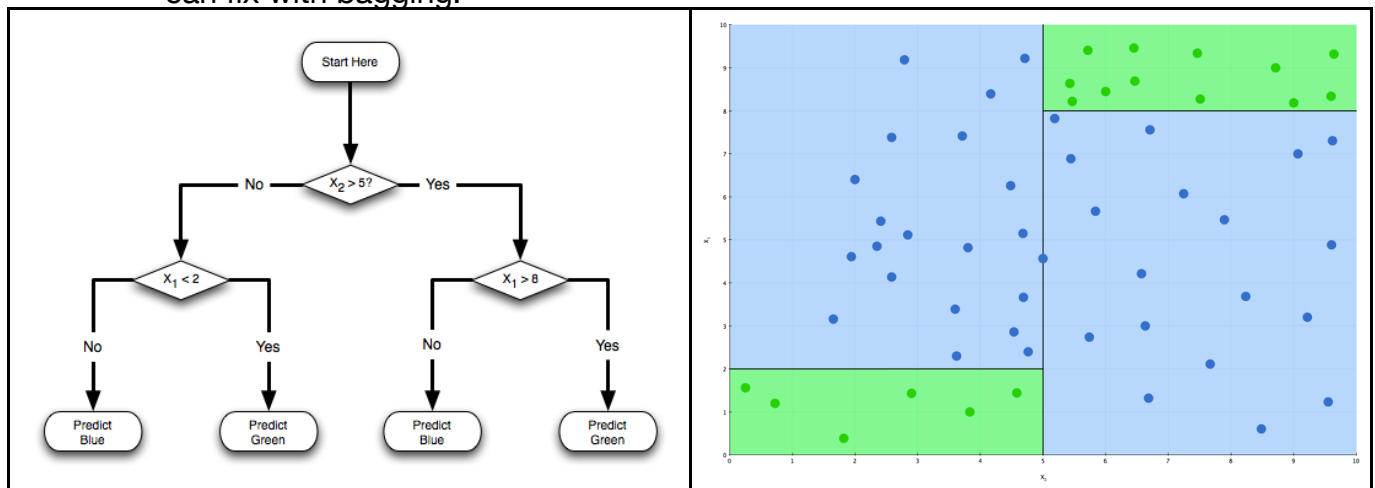
- **Ensemble Classifier:**
 - Combine the outputs of many "weak" classifiers to make a "strong" classifier
- **Training:**
 - In each iteration,
 - training data is re-weighted based on whether it is correctly classified or not.
 - weak classifier focuses on misclassified data from previous iterations.
 - Use cross-validation to pick number of weak learners.
- **Advantages:**
 - Good generalization performance
 - Built-in features selection - decision stump selects one feature at a time.
- **Disadvantages:**
 - Sensitive to outliers.

Outline

1. Nonlinear classifiers
2. Kernel trick and kernel SVM
3. **Ensemble Methods - Boosting, Random Forests**
4. Classification Summary


Decision Tree

- Simple "Rule-based" classifier
 - At each node, move down the tree based on that node's criteria.
 - leaf node contains the prediction
- **Advantage:** can create complex conjunction of rules
- **Disadvantage:** easy to overfit by itself
 - can fix with bagging!



Random Forest Classifier

- Use **bagging** to make an ensemble of Decision Tree Classifiers
 - for each *Decision Tree Classifier*
 - create a new training set by randomly sampling from the training set
 - for each split in a tree, select a random subset of features to use
- for a test sample, the prediction is aggregated over all trees.




```
In [26]: # learn a RF classifier
# use 4 trees
clf = ensemble.RandomForestClassifier(n_estimators=4, random_state=4487, n_jobs=-1)
clf.fit(X3, Y3)
```

```
Out[26]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=4, n_jobs=-1,
                                oob_score=False, random_state=4487, verbose=0,
                                warm_start=False)
```

- Here are the 4 decision trees
 - each uses a different random sampling of original training set

```
In [29]: dtfig
```


```
Out[29]:
```



- and the aggregated classifier

```
In [30]: rffig
```

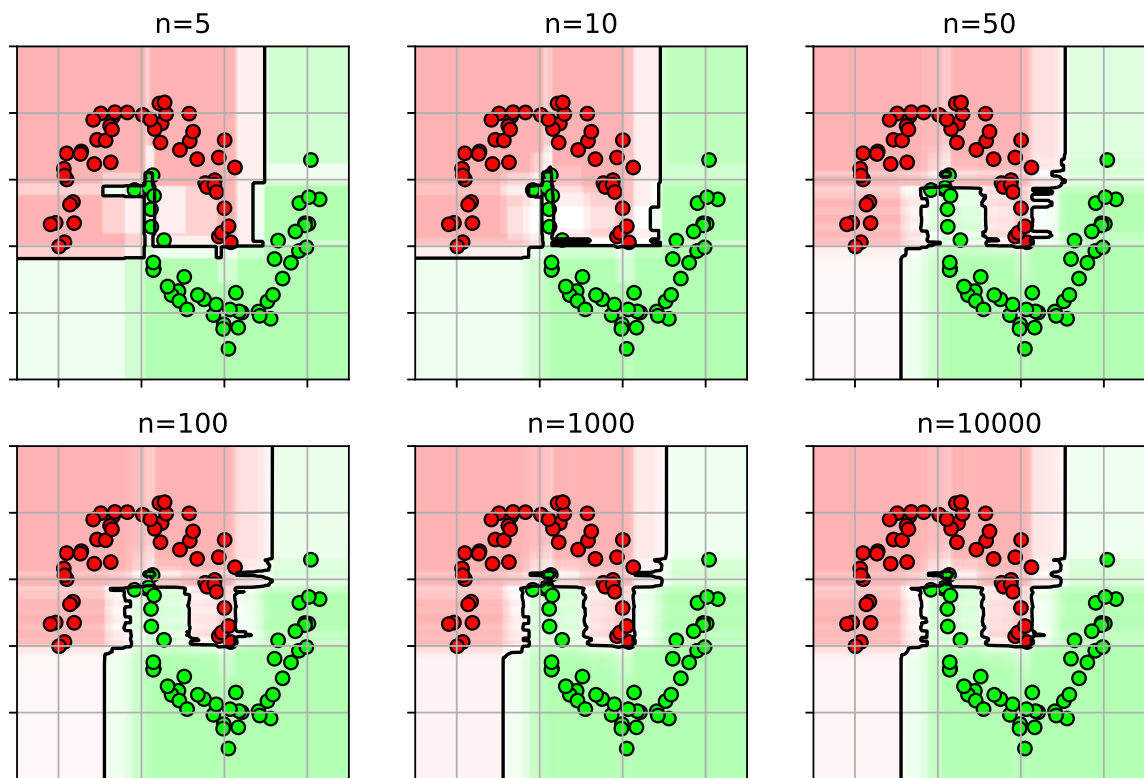
```
Out[30]:
```



- Using more trees

```
In [31]: # learn RF classifiers for different n_estimators
plt.figure(figsize=(9,6))
clfs = {}
for i,n in enumerate([5, 10, 50, 100, 1000, 10000]):
    clfs[n] = ensemble.RandomForestClassifier(n_estimators=n, random_state=4487,
n_jobs=-1)
    clfs[n].fit(X3, Y3)

    plt.subplot(2,3,i+1)
    plot_rf(clfs[n], axbox, X3)
    plt.scatter(X3[:,0], X3[:,1], c=Y3, cmap=mycmap, edgecolors='k')
    plt.gca().xaxis.set_ticklabels([])
    plt.gca().yaxis.set_ticklabels([])
    plt.title("n=" + str(n))
```




- Try on the iris data

```
In [32]: # learn RF classifiers for different n_estimators
plt.figure(figsize=(9,6))
clfs = {}
axbox = [2.5, 7, 1.5, 4]

for i,n in enumerate([5, 10, 50, 100, 500, 1000]):
    clfs[n] = ensemble.RandomForestClassifier(n_estimators=n, random_state=4487,
n_jobs=-1)
    clfs[n].fit(trainX, trainY)

    plt.subplot(2,3,i+1)
    plot_rf(clfs[n], axbox, trainX)
    plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, edgecolors='k')
    plt.gca().xaxis.set_ticklabels([])
    plt.gca().yaxis.set_ticklabels([])
    plt.title("n=" + str(n))
```



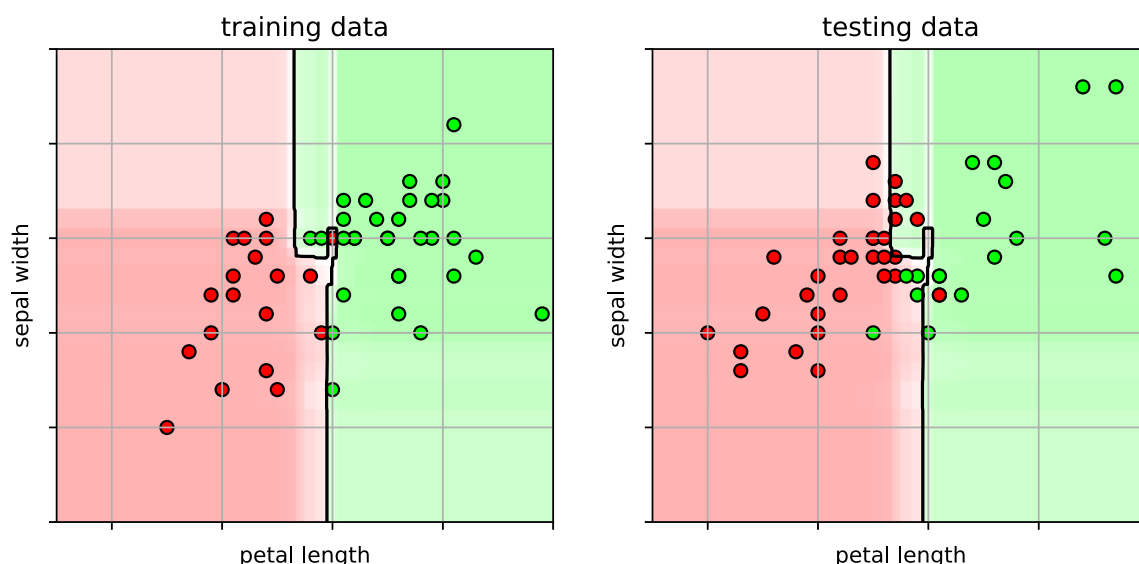
```
In [33]: # predict from the model
predY = clfs[1000].predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)
```

test accuracy = 0.8

```
In [35]: # classifier boundary w/ training and test data
         ifig3
```

Out[35]:



- Important parameters for cross-validation
 - `max_features` - maximum number of features used for each split
 - `max_depth` - maximum depth of a decision tree

Outline

1. Nonlinear classifiers
 2. Kernel trick and kernel SVM
 3. Ensemble Methods - Boosting, Random Forests
 4. **Classification Summary**
-


Feature Pre-processing

- Some classifiers, such as SVM and LR, are sensitive to the scale of the feature values.
 - feature dimensions with larger values may dominate the objective function.
- Common practice is to *standardize* or *normalize* each feature dimension before learning the classifier.
 - Two Methods...
- **Method 1:** scale each feature dimension so the mean is 0 and variance is 1.
 - $\tilde{x}_d = \frac{1}{s}(x_d - m)$
 - s is the standard deviation of feature values.
 - m is the mean of the feature values.
- **NOTE:** the parameters for scaling the features should be estimated from the training set!
 - same scaling is applied to the test set.

```
In [36]: # using the iris data
scaler = preprocessing.StandardScaler() # make scaling object
trainXn = scaler.fit_transform(trainX) # use training data to fit scaling parameters
testXn = scaler.transform(testX) # apply scaling to test data
```

```
In [38]: nfig1
```

Out[38]:




- **Method 2:** scale features to a fixed range, -1 to 1.
 - $\tilde{x}_d = 2 * (x_d - \min) / (\max - \min) - 1$
 - \max and \min are the maximum and minimum features values.

```
In [39]: # using the iris data
scaler = preprocessing.MinMaxScaler(feature_range=(-1,1)) # make scaling object
trainXn = scaler.fit_transform(trainX) # use training data to fit scaling parameters
testXn = scaler.transform(testX) # apply scaling to test data
```



```
In [41]: nfig2
```

```
Out[41]:
```



Data Representation and Feature Engineering

- How to represent data as a vector of numbers?
 - the encoding of the data into a feature vector should make sense
 - inner-products or distances calculated between feature vectors should be meaningful in terms of the data.
- Categorical variables
 - Example: x has 3 possible category labels: cat, dog, horse
 - We could encode this as: $x = 0$, $x = 1$, and $x = 2$.
 - Suppose we have two data points: $x = \text{cat}$, $x' = \text{horse}$.
 - What is the meaning of $x * x' = 2$?

One-hot encoding

- encode a categorical variable as a vector of ones and zeros
 - if there are K categories, then the vector is K dimensions.
- Example:
 - $x = \text{cat} \rightarrow x = [1 \ 0 \ 0]$
 - $x = \text{dog} \rightarrow x = [0 \ 1 \ 0]$
 - $x = \text{horse} \rightarrow x = [0 \ 0 \ 1]$

```
In [42]: # one-hot encoding example
X = [[0], [1], [0], [2], [2]] # original categorical data {0,1,2}
ohe = preprocessing.OneHotEncoder(sparse=False)
ohe.fit(X) # finds the number of categories in the training set: 0-max(X)
ohe.transform(X) # transform to one-hot-encoding
```

```
Out[42]: array([[1., 0., 0.],
 [0., 1., 0.],
 [1., 0., 0.],
 [0., 0., 1.],
 [0., 0., 1.]])
```

Binning

- encode a real value as a vector of ones and zeros
 - assign each feature value to a bin, and then use one-hot-encoding

```
In [43]: # example
X = [[-3], [0.5], [1.5]] # the data
bins = [-2,-1,0,1,2] # define the bins

# map from value to bin number
Xbins = digitize(X, bins=bins)

# map from bin number to 0-1 vector
ohe = preprocessing.OneHotEncoder(n_values=len(bins), sparse=False)
ohe.fit(Xbins)
ohe.transform(Xbins)
```

```
Out[43]: array([[1., 0., 0., 0., 0.],
 [0., 0., 0., 1., 0.],
 [0., 0., 0., 0., 1.]])
```

Data transformations - polynomials

- Represent interactions between features using polynomials
- Example:
 - 2nd-degree polynomial models pair-wise interactions
 - $[x_1, x_2] \rightarrow [x_1^2, x_1x_2, x_2^2]$
 - Combine with other degrees:
 - $[x_1, x_2] \rightarrow [1, x_1, x_2, x_1^2, x_1x_2, x_2^2]$

```
In [44]: X = [[0,1], [1,2], [3,4]]
pf = preprocessing.PolynomialFeatures(degree=2)
pf.fit(X)
pf.transform(X)
```

```
Out[44]: array([[ 1.,  0.,  1.,  0.,  0.,  1.],
 [ 1.,  1.,  2.,  1.,  2.,  4.],
 [ 1.,  3.,  4.,  9., 12., 16.]])
```

Data transformations - univariate


- Apply a non-linear transformation to the feature
 - e.g., $x \rightarrow \log(x)$
 - useful if the dynamic range of x is very large

Unbalanced Data

- For some classification tasks that data will be unbalanced
 - many more examples in one class than the other.
- **Example:** detecting credit card fraud
 - credit card fraud is rare
 - 50 examples of fraud, 5000 examples of legitimate transactions.

```
In [46]: udatafig
```

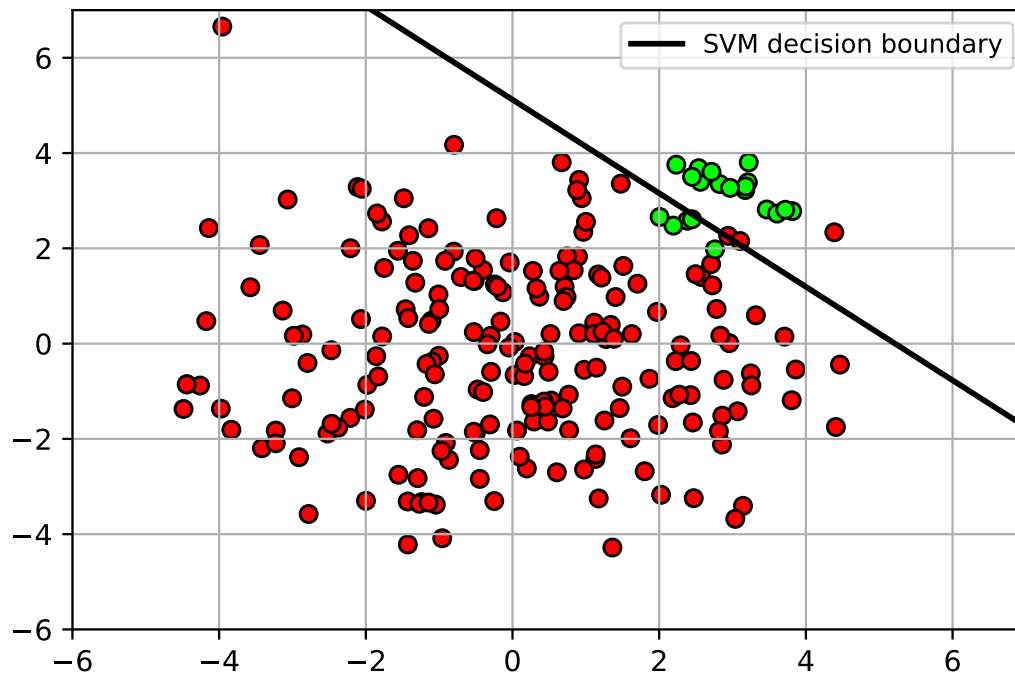
```
Out[46]:
```



- Unbalanced data can cause problems when training the classifier
 - classifier will focus more on the class with more points.
 - decision boundary is pushed away from class with more points

```
In [48]: udatafig1
```

```
Out[48]:
```



- **Solution:** apply weights on the classes during training.
 - weights are inversely proportional to the class size.


```
In [49]: clfw = svm.SVC(kernel='linear', C=10, class_weight='balanced')
         clfw.fit(X, Y)

         print("class weights =", clfw.class_weight_)

class weights = [0.55 5.5 ]
```

```
In [51]: udatafig2
```

```
Out[51]:
```




Classifier Imbalance

- In some tasks, errors on certain classes cannot be tolerated.
- **Example:** detecting spam vs non-spam
 - non-spam should *definitely not* be marked as spam
 - okay to mark some spam as non-spam

```
In [53]: udatafig3
```

```
Out[53]:
```




- Class weighting can be used to make the classifier focus on certain classes
 - e.g., weight non-spam class higher than spam class
 - classifier will try to correctly classify all non-spam samples, at the expense of making errors on spam samples.

```
In [54]: # dictionary (key,value) = (class name, class weight)
cw = {0: 0.2,
      1: 5} # class 1 is 25 times more important!

clfw = svm.SVC(kernel='linear', C=10, class_weight=cw)
clfw.fit(X, Y);
```

```
In [56]: udatafig4
```

Out[56]:



Classification Summary

- **Classification task**

- Observation \mathbf{x} : typically a real vector of feature values, $\mathbf{x} \in \mathbb{R}^d$.
- Class y : from a set of possible classes, e.g., $\mathcal{Y} = \{0, 1\}$
- **Goal:** given an observation \mathbf{x} , predict its class y .

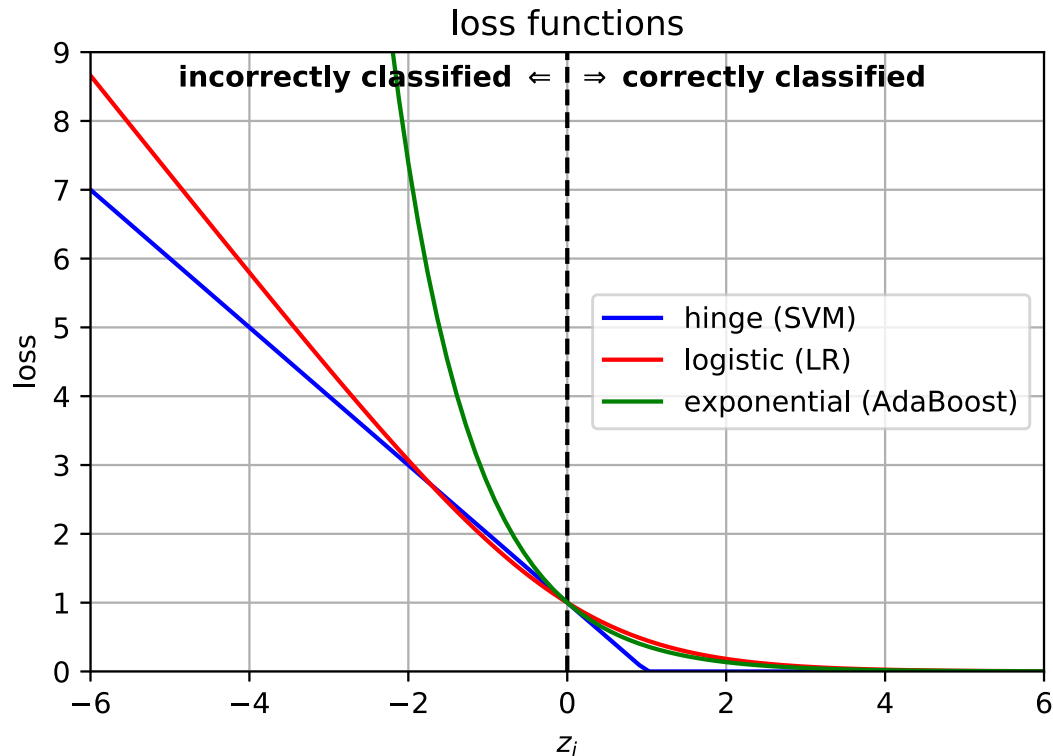
Name	Type	Classes	Decision function	Training	Advantages	Disadvantages
Bayes' classifier	generative	multi-class	non-linear	estimate class-conditional densities $p(\mathbf{x} y)$ by maximizing likelihood of data.	<ul style="list-style-type: none"> - works well with small amounts of data. - multi-class. - minimum probability of error if probability models are correct. 	<ul style="list-style-type: none"> - depends on the data correctly fitting the class-conditional.
logistic regression	discriminative	binary	linear	maximize likelihood of data in $p(y \mathbf{x})$.	<ul style="list-style-type: none"> - well-calibrated probabilities. - efficient to learn. 	<ul style="list-style-type: none"> - linear decision boundary. - sensitive to C parameter.
support vector machine (SVM)	discriminative	binary	linear	maximize the margin (distance between decision surface and closest point).	<ul style="list-style-type: none"> - works well in high-dimension. - good generalization. 	<ul style="list-style-type: none"> - linear decision boundary. - sensitive to C parameter.
kernel SVM	discriminative	binary	non-linear (kernel function)	maximize the margin.	<ul style="list-style-type: none"> - non-linear decision boundary. - can be applied to non-vector data using appropriate kernel. 	<ul style="list-style-type: none"> - sensitive to kernel function and hyperparameters. - high memory usage for large datasets
AdaBoost	discriminative	binary	non-linear (ensemble of weak learners)	train successive weak learners to focus on misclassified points.	<ul style="list-style-type: none"> - non-linear decision boundary. can do feature selection. - good generalization. 	<ul style="list-style-type: none"> - sensitive to outliers.
Random Forest	discriminative	multi-class	non-linear (ensemble of decision trees)	aggregate predictions over several decision trees, trained using different subsets of data.	<ul style="list-style-type: none"> - non-linear decision boundary. can do feature selection. - good generalization. - fast 	<ul style="list-style-type: none"> - sensitive to outliers.

Loss functions

- The classifiers differ in their loss functions, which influence how they work.
 - $z_i = y_i f(\mathbf{x}_i)$

In [58]: `lossfig`

Out[58]:




Regularization and Overfitting

- Some models have terms to prevent overfitting the training data.
 - this can improve *generalization* to new data.
- There is a parameter to control the regularization effect.
 - select this parameter using cross-validation on the training set.


```
In [60]: ofig
```

```
Out[60]:
```



Other things

- *Multiclass classification*
 - can use binary classifiers to do multi-class using *1-vs-rest* formulation.
- *Feature normalization*
 - normalize each feature dimension so that some feature dimensions with larger ranges do not dominate the optimization process.
- *Unbalanced data*
 - if more data in one class, then apply weights to each class to balance objectives.
- *Class imbalance*
 - mistakes on some classes are more critical.
 - reweight class to focus classifier on correctly predicting one class at the expense of others.

Applications

- Web document classification, spam classification
- Face gender recognition, face detection, digit classification

Features

- Choice of features is important!
 - using uninformative features may confuse the classifier.
 - use domain knowledge to pick the best features to extract from the data.

Which classifier is best?

- **"No Free Lunch" Theorem** (Wolpert and Macready)

"If an algorithm performs well on a certain class of problems then it necessarily pays for that with degraded performance on the set of all remaining problems."

- In other words, there is no *best* classifier for all tasks. The best classifier depends on the particular problem.